

[Home](#) > [Posts](#) > [OpenAI drops a security scanner and new voice APIs](#)

OPENAI DROPS A SECURITY SCANNER AND NEW VOICE APIS

Also: Build your local AI coding setup

Jul 29, 2026



THE CODE

Together with **ORACLE**


Welcome back. The recent GPT-5.6 breach into Hugging Face's infrastructure was a major wake up call for every engineering team in the industry. Today, we'll explore how that exploit happened and how OpenAI is patching the holes to stop it from happening again.

Today's Brief

- A developer's guide on spec-driven development
- 10 AI projects that get you hired
- 6 GitHub settings to secure your project
- Oracle's [guide](#) on Agent Loop

TODAY IN PROGRAMMING





Click here to see OpenAI's voice models in action.

OpenAI drops a security scanner and new voice APIs: The ChatGPT maker quietly open-sourced [Codex Security](#), a CLI built on its Codex agent that scans repos for vulnerabilities and suggests patches for developer review. You can install the CLI with a single command: `npm install @openai/codex-security`. The company also shipped APIs for GPT-Live-Transcribe and GPT-Transcribe. The model cuts GPT Whisper's error rate by more than half on real-world benchmarks. More details on the API [here](#).

Hugging Face released an interactive replay of the OpenAI breach: Earlier this month, an autonomous agent powered by OpenAI models broke into Hugging Face during benchmark testing. The model-hosting platform has now published a technical timeline of the incident. The biggest surprise was just how deep and persistent the intrusion from OpenAI's model became. Over 17,600 actions later, it rebuilt its own tooling every time its environment was wiped with no human involved. See exactly [what happened here](#).

Grok drops working prototypes on engineering teams: Elon Musk's AI lab just rolled out [Build Mode](#) in early beta for SuperGrok Heavy subscribers. Instead of just spitting out code snippets to paste elsewhere, it renders a live, working preview right in your chat thread. Once you're done, you can publish projects to a grok.me link or a custom domain. You can even export the source to GitHub.

PRESENTED BY ORACLE

[Decoding the Agent Loop](#)

Agent Harness

Model

reasoning and decision making

context assembly | tool execution | operational constraints | state persistence

Most agent engineering happens in the harness and not the model, making it the exact place where most failures arise. Building successful agent systems depends on understanding the agent loop across its 3 distinct levels.

As these loops scale and evolve, so does the risk of production failure. What unites them is the [memory layer](#), turning stateless loops into a state-aware reasoning engine.

[Access the full agent loop breakdown.](#)

INSIGHT

AI coding falls apart in large codebases. Spec-driven development is the fix:



Source: *The Code, Superhuman*

Guesses compound. For years, teams shipped AI code the same way. You prompt the agent, it fills the gaps with a guess, and you clean up the rest. That worked for small tasks. But it broke down as codebases grew, because the bigger the system, the more bugs turn up in details you never touched. Engineering leader Gregor Ojstersek recently broke down the [workflow](#) that AI-native teams use to lock those details down before the first line of code is ever written.

The fix? Spec-driven development. This means writing a "spec" (a clear description of what you're building, including features, requirements, and constraints) before any AI code is generated. On a large codebase, a good spec has six parts:

- **Problem statement at the top:** Explain what you're solving and for whom.
- **Non-goals:** Write what this system explicitly will not do.
- **Assumptions:** List your decisions so the model doesn't guess them.
- **Reference implementation:** Call out which parts are shortcuts and which parts prove the implementation.
- **Architecture:** Data model, module boundaries, interfaces, error handling, etc.
- **Test plan:** What unit, integration, and/or e2e tests are needed and how to do them.

If you want to see one in practice, here's a [well-written spec](#) for reference.

Not always worth it. A detailed spec is worth the effort for a massive codebase. But for smaller projects, you'll likely spend more time writing the document than actually shipping the feature. If

P.S. You can also watch this tutorial from an Oracle engineer who demonstrates spec-driven development using a simple [markdown file](#).

IN THE KNOW

What's trending on socials and headlines



atlas



@creatine_cycle



dudes will be skinny and be concerned with open source vs closed source models.

bro you're worried about the wrong weights

4:36 AM · Jul 28, 2026 · **245.5K** Views



130



387



8K




484



Meme of the day.

- **Hiring Standouts:** Most devs stay stuck watching tutorials. This thread lays out [10 AI projects](#) that get you hired.
- **Security First:** You don't need to be a security engineer to improve your project's security. Here are [6 GitHub settings](#) you need to enable to make it secure (2.8K bookmarks).
- **Rent a MacBook:** Apple just made it a lot easier to keep your machine spec'd out for running local AI. You can now [lease a MacBook](#) monthly at a way lower base price (1M views).
- **AI Detective:** This developer lost \$590 in API credits in 48 hours. So he ran [one prompt](#) on GLM 5.2 and found the attacker in under 20 minutes (2K likes).
- **Instant Memory:** Agent memory is still a huge problem. This founder lays out the [exact playbook](#) to build a company brain (2.3K bookmarks).

TOP & TRENDING RESOURCES



[Click here to watch the tutorial.](#)

Top Tutorial

[The complete local AI coding setup using Ollama and OpenCode:](#) This tutorial shows you how to set up a local AI coding assistant using Ollama, OpenCode, and VS Code. It walks through configuring local LLMs like Qwen to build Next.js apps, explain React code, and run Jupyter notebooks directly on your local GPU without relying on cloud services.

Top Repo

[no-mistakes \(7.2K 🌟\):](#) This GitHub repo from a Microsoft engineer sets up a local git proxy for your remote. Instead of pushing to 'origin,' you push to 'no-mistakes.' This spins up a temporary worktree and runs an AI validation pipeline. Your branch only moves to the target if every check passes and even opens a clean PR automatically.

Trending Cookbook

[How to review AI generated code \(by a Google Engineer\):](#) AI agents now write code faster than humans can read it, shifting the bottleneck from production to verification. This guide from a former Google engineering leader explains how to adapt. To keep up, developers need to scale their reviews based on risk, using AI for initial triage while saving human oversight for the most critical merges.

AI CODING HACK

Sessions hit context limits early. Too much space is wasted before the first prompt on MCP servers, unused skills, and an ever-expanding CLAUDE.md. Claude Code v2.1.206 introduced an audit to fix this. It identifies the bloat and proposes removals for you to approve.

- **Step 1:** Run “/context” and record the total.

```
# Show what occupies the window before the first prompt
/context
```

```
# Audit skills, MCP servers, plugins, hooks, and CLAUDE.md
/doctor
```

```
# Compare against the starting number
/context
```

- **Step 2:** Run “/doctor”. It weighs skills, MCP servers, and plugins against their context cost and identifies CLAUDE.md content Claude can derive from the codebase.
- **Step 3:** Approve the changes. Remove directory layouts, dependency lists, and build commands. Keep the pitfalls and conventions.

P.S. Get 50+ AI coding hacks for Claude Code, Cursor, and Codex [here](#).

IN CASE YOU MISSED IT

Our most-clicked story from yesterday

Find out how you can run Moonshot AI's [Kimi K3 locally](#).

Grow customers & revenue: Join companies like Google, IBM, and Datadog. Showcase your product to our 300K+ engineers and 150K+ followers on socials. [Get in touch](#).

What did you think of today's newsletter?

Your feedback helps us create better emails for you!

Loved it 🧠🧠🧠


Average 🧠🧠

It sucked 🧠

[Login](#) or [Subscribe](#) to participate

Until next time — The Code team


KEEP READING



Aug 24, 2026 • 5 min read

VERCEL SHIPS TOOL THAT GRADES SITES FOR AI ...

Also: How to build an autonomous AI coding system



Aug 21, 2026 • 4 min read

A MYSTERIOUS MODEL IS GOING VIRAL AMONG ...

Also: A senior engineer's complete development system



Sep 2, 2026 • 4 min read

ANTHROPIC MAKES ITS BEST CODING MODEL ...

Also: Build an AI Software Factory

View more >

Email

Subscribe

Archive

LinkedIn

Superhuman

Advertise

Marketing
Standup